



InfiniBand Architecture

High-Performance Interconnect for HPC & Data Centers

DAC STRATEGIC ELECTRONICS GROUP

Introduction to InfiniBand

Next-Generation Interconnect Standard

InfiniBand (IB) is a high-performance network communications standard designed specifically for high-performance computing (HPC), AI/ML workloads, and enterprise networking. It overcomes traditional Ethernet limitations to provide a critical foundation for modern data-intensive applications.

Key Applications

- **HPC Supercomputers** & AI/ML training clusters
- **High-Speed Storage** networks and database fabrics
- **Enterprise Data Centers** & modern cloud infrastructure

Advanced Capabilities

Remote Direct Memory Access (RDMA) enables zero-copy data transfers, bypassing CPU overhead completely.

Quality of Service (QoS) guarantees bandwidth and dedicated virtual lanes under heavy computational loads.



Bandwidth ULTRA-HIGH

Provides massive data throughput to prevent bottlenecks during intensive parallel computations.

Latency MICROSECOND

Minimizes transit delays for hardware-based, real-time node-to-node communications.

Scalability LINEAR

Allows seamless expansion to tens of thousands of active nodes across complex topologies.

Key Features of InfiniBand

High Bandwidth Evolution

From SDR (Single Data Rate) (2.5 Gbps) to current NDR (Next Data Rate) (400 Gbps), InfiniBand has consistently pushed bandwidth boundaries. Each generation delivers 4x improvement, with HDR (High Data Rate) (200 Gbps) currently dominating HPC installations and NDR (Next Data Rate) systems emerging for exascale computing.

Ultra-Low Latency

Nanosecond-level switching latency enables real-time communication between compute nodes. Hardware-based switching and cut-through forwarding minimize packet processing delays, critical for tightly-coupled parallel applications.

RDMA Support

Remote Direct Memory Access allows zero-copy data transfers, bypassing CPU involvement and eliminating kernel overhead. This dramatically reduces latency and frees CPU cycles for computational tasks rather than data movement.

Massive Scalability


InfiniBand fabrics scale to tens of thousands of nodes through hierarchical topologies like fat-tree and dragonfly. Advanced routing algorithms ensure optimal path selection and load balancing across the entire fabric.

Quality of Service

Virtual Lanes (VLs) provide traffic isolation and priority handling, ensuring critical communications receive guaranteed bandwidth and latency bounds even under heavy network load.


Massive Scalability

InfiniBand fabrics scale to tens of thousands of nodes through hierarchical topologies like fat-tree and dragonfly. Advanced routing algorithms ensure optimal path selection and load balancing across the entire fabric.



Quality of Service

Virtual Lanes (VLs) provide traffic isolation and priority handling, ensuring critical communications receive guaranteed bandwidth and latency bounds even under heavy network load.



InfiniBand Components



Channel Adapters

Host Channel Adapters (HCA)

connect servers and compute nodes to the InfiniBand fabric, providing hardware-accelerated networking with RDMA capabilities. **Target Channel Adapters (TCA)** connect storage devices and I/O systems, enabling high-speed data access with minimal CPU overhead.



Switches & Routers

InfiniBand switches provide low-latency packet forwarding within subnets, supporting up to 36 ports with line-rate performance. Routers enable inter-subnet communication and larger fabric scaling, maintaining performance across subnet boundaries.



Subnet Manager

The Subnet Manager (SM) serves as the fabric's control plane, discovering topology, assigning addresses, configuring routing tables, and monitoring fabric health. It ensures optimal performance and automatic recovery from failures.

These components work together to create a unified, high-performance fabric that appears as a single, coherent network to applications. The hardware-accelerated design offloads network processing from the CPU, enabling maximum computational efficiency.

InfiniBand Communication Model

Queue Pairs (QP) Architecture

InfiniBand communication revolves around Queue Pairs - the fundamental communication endpoint consisting of a Send Queue and Receive Queue. Each application creates QPs to establish communication channels with remote applications, enabling direct memory-to-memory transfers.

Work Requests & Elements

Applications post Work Requests (WR) to queues, which are converted to Work Queue Elements (WQE) by the hardware. These elements contain all information needed for data transfer, including memory addresses, transfer sizes, and operation types.

Completion Notifications

Completion Queues (CQ) provide asynchronous notification when operations complete, allowing applications to overlap computation with communication for maximum efficiency.



01

Application Posts Work Request

Application submits send or receive operation to appropriate queue

02

Hardware Processes WQE

Channel adapter converts request to work queue element and initiates transfer

03

Data Transfer Occurs

Direct memory access between source and destination without CPU involvement

04

Completion Generated

Hardware posts completion entry to completion queue for application notification

Key Advantage: The QP model enables zero-copy transfers and CPU bypass, dramatically reducing communication overhead compared to traditional socket-based networking.

Transport Services

InfiniBand provides multiple transport services to match different application requirements, balancing performance, reliability, and resource utilization based on specific use cases.



Reliable Connection (RC)

Point-to-point reliable service with guaranteed delivery and ordering.

Provides end-to-end flow control and automatic retransmission.

Ideal for critical data transfers, commonly used in storage and database applications.



Unreliable Connection (UC)

Point-to-point service without hardware reliability guarantees.

Offers lower overhead and higher performance; errors are handled by applications.

Suitable for streaming media or applications with built-in recovery.



Reliable Datagram (RD)

Connectionless reliable service supporting multicast communication.

Provides guaranteed delivery with automatic retransmission.

Requires more complex addressing, used for group communications.



Unreliable Datagram (UD)

Connectionless service with minimal overhead, supporting unicast and multicast.

No reliability guarantees but offers maximum performance and scalability.

Commonly used for management traffic and discovery protocols.



Raw Datagram Service

Allows transmission of non-InfiniBand protocols over the InfiniBand fabric, enabling legacy application support and protocol tunneling. This service maintains the high-performance characteristics of InfiniBand while supporting existing network protocols and applications.

C-DAC STRATEGIC ELECTRONICS GROUP

InfiniBand Addressing

Local Identifier (LID)

16-bit addresses used within a single subnet, supporting up to 65,536 unique addresses. LIDs are assigned by the Subnet Manager and enable efficient intra-subnet routing. Each port receives a base LID, with additional LIDs available for multipath configurations to increase bandwidth and provide redundancy. LID-based routing uses simple forwarding tables in switches, enabling nanosecond-level switching decisions. The hierarchical addressing scheme allows for efficient fabric utilization and optimal path selection within the subnet boundaries.

Global Identifier (GID)

128-bit IPv6-compatible addresses enabling communication across subnet boundaries. GIDs consist of a 64-bit subnet prefix and 64-bit interface identifier, providing globally unique addressing for large-scale deployments spanning multiple subnets. GID-based routing enables complex multi-subnet topologies and integration with IP networks. This addressing scheme supports the massive scale requirements of modern supercomputing installations and cloud deployments.



Subnet Manager Role

Discovers fabric topology, assigns LIDs to all ports, calculates optimal routing paths, and maintains forwarding tables across all switches.

Address Assignment


Dynamically allocates addresses based on fabric topology, ensures no conflicts, and optimizes routing for maximum performance and fault tolerance.

Path Optimization

Calculates multiple paths between endpoints, balances traffic load, and provides automatic failover capabilities for high availability.

InfiniBand Packet Structure

InfiniBand packets are carefully structured to enable high-performance, reliable communication while maintaining compatibility across different network topologies and transport services.



Virtual Lanes & Flow Control


Virtual Lanes (VLs) Architecture

InfiniBand supports up to 16 Virtual Lanes (VL0-VL15) that provide traffic isolation and quality of service within a single physical link. Each VL maintains separate buffer space and scheduling priority, preventing Head-of-Line (HoL) blocking where high-priority traffic can be delayed by lower-priority packets.

VLs enable sophisticated traffic management policies, allowing administrators to dedicate specific lanes to critical applications, management traffic, or different user groups. This granular control ensures predictable performance even under heavy network load conditions.


Credit-Based Flow Control

InfiniBand implements a sophisticated credit-based flow control mechanism where receivers advertise available buffer space to senders. Each VL maintains separate credit pools, ensuring that congestion in one lane doesn't affect others. This prevents packet loss and maintains optimal throughput across all traffic classes.



Maximum VLs

Up to 16 virtual lanes per physical link for traffic isolation



Packet Loss

Credit-based flow control eliminates packet drops due to congestion



Priority Levels

Configurable priority scheduling across virtual lanes



Performance Benefit: VLs eliminate head-of-line blocking, ensuring that high-priority traffic maintains low latency even when lower-priority traffic saturates the network.

Future of InfiniBand



NDR & XDR Evolution

NDR (400 Gbps) systems are entering production for exascale supercomputers, while XDR (800 Gbps) development targets future extreme-scale computing requirements. These advances maintain InfiniBand's leadership in HPC interconnects.

Market Leadership & Innovation

InfiniBand continues to dominate the HPC interconnect market, powering over 60% of the world's fastest supercomputers. Its roadmap includes advanced features like in-network computing, enhanced security, and tighter integration with emerging computing paradigms including quantum and neuromorphic systems.

The technology's evolution focuses on maintaining its performance leadership while expanding into new domains such as edge computing, autonomous vehicles, and real-time analytics where ultra-low latency and high reliability are paramount.



Exascale Computing

InfiniBand provides the foundational interconnect for exascale systems like Frontier and Aurora, enabling unprecedented computational capabilities through ultra-low latency and massive bandwidth scaling.



AI/ML Integration

Deep integration with GPU clusters through CUDA-aware MPI and GPUDirect technologies enables direct GPU-to-GPU communication, accelerating AI training and inference workloads across distributed systems.



"InfiniBand remains the gold standard for high-performance interconnects, continuously pushing the boundaries of what's possible in distributed computing and data-intensive applications."

C-DAC STRATEGIC ELECTRONICS GROUP

Intel Omni-Path Architecture & Packet Structure

High-Performance Computing Interconnect Technology

C-DAC STRATEGIC ELECTRONICS GROUP



Introduction to Omni-Path Architecture

Intel Omni-Path Architecture (OPA) represents a revolutionary high-performance interconnect technology designed specifically for demanding HPC environments. Developed to overcome the scalability limitations of traditional InfiniBand solutions, OPA addresses the growing computational demands of modern supercomputing applications.

OPA was engineered to support massive-scale deployments where traditional interconnects fall short. The architecture targets critical use cases including large-scale HPC clusters, artificial intelligence and machine learning workloads, and next-generation supercomputing systems that require unprecedented levels of performance and scalability.

The technology emerged from Intel's recognition that existing interconnect solutions couldn't meet the exponential growth in computational requirements, particularly for applications requiring low-latency, high-bandwidth communication across thousands of nodes simultaneously.



Key Innovation: OPA was specifically designed to scale beyond 100,000 nodes while maintaining deterministic performance characteristics.



Design Goals of OPA

Massive Scalability

Support for **>100,000 nodes** in a single fabric, enabling unprecedented cluster sizes for next-generation supercomputing applications.

High Bandwidth

100 Gbps per port with scalable architecture supporting multiple ports per node for aggregate bandwidth exceeding terabits per second.

Ultra-Low Latency


Approximately **~100 ns per hop** latency, critical for tightly-coupled parallel applications requiring frequent inter-node communication.

High Port Density

48-64 ports per switch ASIC, significantly higher than competing solutions, reducing infrastructure complexity and cost.

These design goals work synergistically to deliver deterministic congestion control and adaptive routing capabilities, ensuring consistent performance even under heavy network loads typical in large-scale HPC environments.

Components of OPA



Host Fabric Interface (HFI)

The HFI serves as the network interface card, providing essential NIC functions while offering advanced RDMA offload capabilities. It handles protocol processing, memory management, and direct memory access operations, significantly reducing CPU overhead during network operations. Key features include zero-copy data transfers, hardware-accelerated RDMA operations, and integrated support for multiple virtual lanes to ensure quality of service across different traffic types.



Switch ASICs

Purpose-built switch ASICs deliver exceptional port density while maintaining power efficiency. These custom silicon solutions provide 48-64 ports per ASIC, dramatically higher than traditional switching solutions. The ASICs incorporate advanced buffering, adaptive routing algorithms, and congestion management features, enabling deterministic performance across large-scale fabric deployments.



Fabric Manager

The centralized Fabric Manager provides comprehensive topology management, intelligent routing optimization, and quality of service enforcement across the entire fabric. It continuously monitors fabric health, automatically reconfigures routing tables for optimal performance, and manages virtual lane assignments to ensure predictable application performance.

OPA vs InfiniBand Comparison

Feature	Intel OPA	InfiniBand
Switch Port Density	48-64 ports	36 ports
Bandwidth per Port	100 Gbps	100 Gbps
Latency per Hop	~100 ns	~100 ns
Maximum Scalability	>100,000 nodes	~10,000 nodes
Congestion Control	Advanced adaptive	Basic
Determinism	Hardware-enforced	Software-dependent

The comparison reveals OPA's significant advantages in scalability and switch density. While both technologies offer similar per-port performance, OPA's superior port density reduces the number of switch tiers required, simplifying fabric topology and reducing latency. The advanced congestion control and hardware-enforced determinism make OPA particularly suitable for large-scale HPC deployments where predictable performance is critical.

OPA Features & Advantages



RDMA Capabilities

Remote Direct Memory Access enables zero-copy data transfers and CPU offload, dramatically reducing latency and improving application performance. Applications can directly access remote memory without kernel intervention.



Adaptive Routing

Dynamic path selection automatically routes packets around congested or failed links, ensuring optimal performance and fault tolerance. The fabric continuously adapts to changing network conditions.



Quality of Service

Virtual Lanes (VLs) provide traffic isolation and prioritization, ensuring critical applications receive guaranteed bandwidth and latency characteristics even in mixed workload environments.



Enhanced Reliability

Comprehensive error detection using CRC checksums combined with automatic retransmission ensures data integrity across the fabric, critical for long-running HPC applications.




Multiple Traffic Classes

Support for diverse traffic types including bulk data transfers, latency-sensitive messages, and management traffic, each with appropriate quality of service guarantees.

C-DAC STRATEGIC ELECTRONICS GROUP

Omni-Path Packet Structure Overview

The OPA packet structure is meticulously designed to support high-performance, reliable communication across large-scale fabrics. Each packet contains multiple headers and integrity checks to ensure proper routing, flow control, and data integrity.



Header Hierarchy

Link Transfer Header (LTH): Fabric-level routing and virtual lane information

Transport Header (TH): End-to-end communication control and sequencing

Payload: Application data, RDMA operations, or management messages

Design Principles

The packet structure balances efficiency with functionality, providing comprehensive routing and reliability features while minimizing overhead. Each header serves specific purposes in the communication process.

Link Transfer Header (LTH) Details

The Link Transfer Header manages fabric-level packet routing and quality of service. It contains essential information for switches to make forwarding decisions and ensure proper traffic handling across the fabric.



Source & Destination LIDs

SLID (Source LID) and DLID (Destination LID) provide fabric addressing, enabling switches to route packets to their intended destinations across complex topologies.



Service Class

Defines the quality of service requirements for the packet, influencing routing decisions and resource allocation throughout the fabric path.



Virtual Lane Selection

VL field determines which virtual lane the packet uses, enabling traffic isolation and preventing head-of-line blocking between different traffic types.



Packet Length

Specifies the total packet size, enabling switches to properly buffer and forward packets while managing flow control credits.



LID Range: OPA supports extended LID ranges to accommodate massive fabric scales exceeding 100,000 nodes.

Transport Header & Payload Structure

Transport Header (TH) Components

The Transport Header manages end-to-end communication reliability and flow control between communicating processes.

OpCodes: Define operation types including Send, RDMA Read/Write, and Atomic operations

Packet Sequence Number (PSN): Ensures proper packet ordering and enables duplicate detection

Flow Control Info: Manages credit-based flow control to prevent buffer overflow

Context ID: Associates packets with specific queue pairs for proper delivery

Payload Section Characteristics

The payload carries the actual application data, RDMA messages, or fabric management information. OPA supports multiple MTU sizes (2 KB, 4 KB, 8 KB) to optimize performance for different application patterns.

2KB

Minimum MTU


Suitable for latency-sensitive
small messages

8KB

Maximum MTU

Optimized for high-bandwidth
bulk transfers

The flexible MTU sizing allows applications to optimize for either latency (smaller packets) or bandwidth efficiency (larger packets) based on their specific communication patterns and performance requirements.




Advanced OPA Features: Virtual Lanes & Reliability

Integrity & Reliability

Dual-level CRC protection with Link CRC (LCRC) for per-hop integrity and End-to-End CRC (ECRC) across the entire fabric path. Automatic retry mechanisms ensure reliable delivery.

Packet Types

Support for diverse packet types including Data Packets, ACK/NACK acknowledgments, Management Packets, and Congestion Notification Packets for comprehensive fabric operation.



Virtual Lanes (VLs)

16 virtual lanes (VL0-VL15) provide traffic isolation and QoS. VL0 typically reserved for management, VL1 for high-priority traffic, and VL2-VL15 for user applications.

Flow Control & Congestion

Credit-based flow control prevents buffer overflow, while advanced congestion detection and adaptive routing ensure fair bandwidth distribution across all fabric users.

These advanced features work together to provide a robust, scalable interconnect solution capable of supporting the most demanding HPC workloads. The combination of hardware-enforced reliability, intelligent traffic management, and adaptive routing makes OPA ideal for next-generation supercomputing applications requiring predictable, high-performance communication across massive node counts.